

# Project: Customer Segmentation using K-Means Clustering

Done By: *Mr. Aridweep Majumder, M.Tech*

This project applies K-Means clustering to a marketing dataset to segment customers based on demographics and purchase behavior.

**Dataset source:** *marketing\_campaign.csv*

**Tools used:** *Python, Pandas, Matplotlib, Seaborn, Scikit-learn*

## ✓ 1. Importing Necessary Libraries & Load Datasets

```
# Basic imports for analysis and modeling
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

# Enable inline plots
%matplotlib inline

# Load the dataset (already uploaded to Colab)
df = pd.read_csv('/content/marketing_campaign.csv')

# Check shape and preview
print(f"Total Rows: {df.shape[0]}")
print(f"Total Columns: {df.shape[1]}")
df.head(3)
```



Total Rows: 2240  
Total Columns: 29

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer	Recency
0	5524	1957	Graduation	Single	58138.0	0	0	04-09-2012	58
1	2174	1954	Graduation	Single	46344.0	1	1	08-03-2014	38
2	4141	1965	Graduation	Together	71613.0	0	0	21-08-2013	26

3 rows × 29 columns

## ✓ 2. Data Overview & Summary

```
# Quick look at column types and names
df.info()
```

```

⇒ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                     2240 non-null   int64
1   Year_Birth             2240 non-null   int64
2   Education              2240 non-null   object
3   Marital_Status         2240 non-null   object
4   Income                 2216 non-null   float64
5   Kidhome                2240 non-null   int64
6   Teenhome               2240 non-null   int64
7   Dt_Customer            2240 non-null   object
8   Recency                2240 non-null   int64
9   MntWines               2240 non-null   int64
10  MntFruits              2240 non-null   int64
11  MntMeatProducts        2240 non-null   int64
12  MntFishProducts        2240 non-null   int64
13  MntSweetProducts       2240 non-null   int64
14  MntGoldProds           2240 non-null   int64
15  NumDealsPurchases      2240 non-null   int64
16  NumWebPurchases        2240 non-null   int64
17  NumCatalogPurchases    2240 non-null   int64
18  NumStorePurchases      2240 non-null   int64
19  NumWebVisitsMonth       2240 non-null   int64
20  AcceptedCmp3           2240 non-null   int64
21  AcceptedCmp4           2240 non-null   int64
22  AcceptedCmp5           2240 non-null   int64
23  AcceptedCmp1           2240 non-null   int64
24  AcceptedCmp2           2240 non-null   int64
25  Complain               2240 non-null   int64
26  Z_CostContact           2240 non-null   int64
27  Z_Revenue              2240 non-null   int64
28  Response               2240 non-null   int64
dtypes: float64(1), int64(25), object(3)
memory usage: 507.6+ KB

```

```

# Listing all columns
print("Columns:")
print(df.columns.tolist())

```

```

⇒ Columns:
['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Income', 'Kidhome', 'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts', 'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth', 'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1', 'AcceptedCmp2', 'Complain', 'Z_CostContact', 'Z_Revenue', 'Response']

```

```

# Check for missing values
print("Missing values in each column:")
print(df.isnull().sum())

```

```

⇒ Missing values in each column:
ID                0
Year_Birth        0
Education         0
Marital_Status    0
Income            24
Kidhome           0
Teenhome          0
Dt_Customer       0
Recency           0
MntWines          0
MntFruits         0
MntMeatProducts   0
MntFishProducts   0
MntSweetProducts  0
MntGoldProds      0
NumDealsPurchases 0
NumWebPurchases   0
NumCatalogPurchases 0
NumStorePurchases 0
NumWebVisitsMonth 0
AcceptedCmp3      0
AcceptedCmp4      0
AcceptedCmp5      0
AcceptedCmp1      0
AcceptedCmp2      0
Complain          0
Z_CostContact     0
Z_Revenue         0
Response          0

```

```

MntSweetProducts      0
MntGoldProds          0
NumDealsPurchases     0
NumWebPurchases       0
NumCatalogPurchases   0
NumStorePurchases     0
NumWebVisitsMonth     0
AcceptedCmp3          0
AcceptedCmp4          0
AcceptedCmp5          0
AcceptedCmp1          0
AcceptedCmp2          0
Complain              0
Z_CostContact         0
Z_Revenue             0
Response              0
dtype: int64

```

```

# Basic statistical summary
df.describe()

```



	ID	Year_Birth	Income	Kidhome	Teenhome	Recency	MntWines	MntI
<b>count</b>	2240.000000	2240.000000	2216.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.
<b>mean</b>	5592.159821	1968.805804	52247.251354	0.444196	0.506250	49.109375	303.935714	26.
<b>std</b>	3246.662198	11.984069	25173.076661	0.538398	0.544538	28.962453	336.597393	39.
<b>min</b>	0.000000	1893.000000	1730.000000	0.000000	0.000000	0.000000	0.000000	0.
<b>25%</b>	2828.250000	1959.000000	35303.000000	0.000000	0.000000	24.000000	23.750000	1.
<b>50%</b>	5458.500000	1970.000000	51381.500000	0.000000	0.000000	49.000000	173.500000	8.
<b>75%</b>	8427.750000	1977.000000	68522.000000	1.000000	1.000000	74.000000	504.250000	33.
<b>max</b>	11191.000000	1996.000000	66666.000000	2.000000	2.000000	99.000000	1493.000000	199.

8 rows × 26 columns

### ✓ 3. Data Cleaning and Preprocessing

```
data = df.copy() # Make a copy to avoid modifying the original dataframe
```

```
data.isnull().sum() # Checking again for missing values
```



	0
<hr/>	
ID	0
Year_Birth	0
Education	0
Marital_Status	0
Income	24
Kidhome	0
Teenhome	0
Dt_Customer	0
Recency	0
MntWines	0
MntFruits	0
MntMeatProducts	0
MntFishProducts	0
MntSweetProducts	0
MntGoldProds	0
NumDealsPurchases	0
NumWebPurchases	0
NumCatalogPurchases	0
NumStorePurchases	0
NumWebVisitsMonth	0
AcceptedCmp3	0
AcceptedCmp4	0
AcceptedCmp5	0
AcceptedCmp1	0
AcceptedCmp2	0
Complain	0
Z_CostContact	0
Z_Revenue	0
Response	0

**dtype:** int64

```
# 'Income' has some missing values - dropping those rows
data = data.dropna(subset=['Income'])
```

```
# Convert 'Dt_Customer' to datetime format
data['Dt_Customer'] = pd.to_datetime(data['Dt_Customer'], format='%d-%m-%Y')
```

```
# Create a new feature: year the customer joined
data['Customer_Year'] = data['Dt_Customer'].dt.year
```

```
# Create a new feature: year the customer joined
data['Customer_Year'] = data['Dt_Customer'].dt.year
```

```
# Drop unused or duplicate columns for clustering
data = data.drop(['ID', 'Dt_Customer'], axis=1)
```

🔗 /tmp/ipython-input-139-3013036707.py:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/](https://pandas.pydata.org/pandas-docs/stable/user_guide/)

```
data['Dt_Customer'] = pd.to_datetime(data['Dt_Customer'], format='%d-%m-%Y')
```

/tmp/ipython-input-139-3013036707.py:8: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/](https://pandas.pydata.org/pandas-docs/stable/user_guide/)

```
data['Customer_Year'] = data['Dt_Customer'].dt.year
```

/tmp/ipython-input-139-3013036707.py:12: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/](https://pandas.pydata.org/pandas-docs/stable/user_guide/)

```
data['Customer_Year'] = data['Dt_Customer'].dt.year
```

```
# Convert 'Education' and 'Marital_Status' to numeric using one-hot encoding
data = pd.get_dummies(data, columns=['Education', 'Marital_Status'], drop_first=True)
```

```
# Convert all boolean columns to integer (so True/False becomes 1/0)
data = data.astype(int)# Final structure check
print(f"Data shape after cleaning: {data.shape}")
data.head(4)
```

🔗 Data shape after cleaning: (2216, 37)

	Year_Birth	Income	Kidhome	Teenhome	Recency	MntWines	MntFruits	MntMeatProducts	MntFishI
0	1957	58138	0	0	58	635	88	546	
1	1954	46344	1	1	38	11	1	6	
2	1965	71613	0	0	26	426	49	127	
3	1984	26646	1	0	26	11	4	20	

4 rows × 37 columns

```
# Final structure check
print(f"Data shape after cleaning: {data.shape}")
data.head()
```

➡ Data shape after cleaning: (2216, 37)

	Year_Birth	Income	Kidhome	Teenhome	Recency	MntWines	MntFruits	MntMeatProducts	MntFish
0	1957	58138	0	0	58	635	88	546	
1	1954	46344	1	1	38	11	1	6	
2	1965	71613	0	0	26	426	49	127	
3	1984	26646	1	0	26	11	4	20	
4	1981	58293	1	0	94	173	43	118	

5 rows × 37 columns

## ✓ 4. Feature Selection for Clustering

```
# Selecting numerical features for clustering
features = data.copy()

# Standardizing the data to bring all features to the same scale
scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)

# Convert back to DataFrame for easier inspection if needed
scaled_df = pd.DataFrame(scaled_features, columns=features.columns)

scaled_df.head(4)
```

➡

	Year_Birth	Income	Kidhome	Teenhome	Recency	MntWines	MntFruits	MntMeatProducts	MntFi
0	-0.986443	0.234063	-0.823039	-0.928972	0.310532	0.978226	1.549429	1.690227	
1	-1.236801	-0.234559	1.039938	0.909066	-0.380509	-0.872024	-0.637328	-0.717986	
2	-0.318822	0.769478	-0.823039	-0.928972	-0.795134	0.358511	0.569159	-0.178368	
3	1.266777	-1.017239	1.039938	-0.928972	-0.795134	-0.872024	-0.561922	-0.655551	

4 rows × 37 columns


## ✓ 5. Finding Optimal Clusters (Elbow Method)

```
# Trying different values of K to see where inertia drops sharply
inertia = []
k_range = range(1, 11)

for k in k_range:
    model = KMeans(n_clusters=k, random_state=42)
    model.fit(scaled_df)
    inertia.append(model.inertia_)

# Plotting the elbow curve
plt.figure(figsize=(8, 5))
```

```
plt.plot(k_range, inertia, marker='o')
plt.title('Elbow Method - Optimal K')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Inertia')
plt.grid(True)
plt.show()
```



## ✓ 6. Apply K-Means Clustering & Visualize Clusters

```
# Let's choose 4 clusters (you can change this based on the elbow plot)
k = 4
```

```
# Fit the model
kmeans = KMeans(n_clusters=k, random_state=42)
clusters = kmeans.fit_predict(scaled_df)
```

```
# Add cluster labels to original data
data['Cluster'] = clusters
```


```
# Visualizing clusters using PCA
```

```
pca = PCA(n_components=2)
pca_data = pca.fit_transform(scaled_df)
```

```
# DataFrame for plotting
pca_df = pd.DataFrame(data=pca_data, columns=['PCA1', 'PCA2'])
pca_df['Cluster'] = clusters
```

```
plt.figure(figsize=(8, 6))
sns.scatterplot(data=pca_df, x='PCA1', y='PCA2', hue='Cluster', palette='tab10', s=70)
```

```
sns.scatterplot(data=pca_df, x='PCA1', y='PCA2', hue='Cluster', palette='dark', s=70,
plt.title('Customer Segments Visualized (PCA)')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.legend(title='Cluster')
plt.grid(True)
plt.show()
```



## ✓ 7. Cluster Analysis

```
# Mean profile of each cluster
cluster_profile = data.groupby('Cluster').mean(numeric_only=True).round(2)
cluster_profile.T # Transposed for easier reading
```





Cluster	0	1	2	3
Year_Birth	1972.44	1969.62	1969.09	1966.21
Income	37261.54	38559.86	81055.89	69113.32
Kidhome	0.71	0.71	0.05	0.07
Teenhome	0.48	0.56	0.15	0.53
Recency	47.41	49.10	48.05	49.79
MntWines	92.84	96.33	863.22	524.26
MntFruits	7.75	5.82	55.62	53.81
MntMeatProducts	38.91	34.30	465.02	318.65
MntFishProducts	9.99	8.91	82.08	75.77
MntSweetProducts	7.36	5.84	64.61	53.59
MntGoldProds	25.09	20.87	77.81	73.78
NumDealsPurchases	2.48	2.55	1.14	2.27
NumWebPurchases	2.88	2.87	5.44	5.85
NumCatalogPurchases	0.94	0.82	6.07	4.97
NumStorePurchases	3.83	3.82	8.23	8.63
NumWebVisitsMonth	6.53	6.44	3.11	3.90
AcceptedCmp3	0.10	0.06	0.15	0.06
AcceptedCmp4	0.02	0.04	0.38	0.06
AcceptedCmp5	0.00	0.00	0.83	0.00
AcceptedCmp1	0.01	0.00	0.49	0.05
AcceptedCmp2	0.01	0.00	0.13	0.00
Complain	0.01	0.01	0.01	0.01
Z_CostContact	3.00	3.00	3.00	3.00
Z_Revenue	11.00	11.00	11.00	11.00
Response	0.16	0.08	0.58	0.12
Customer_Year	2013.08	2013.07	2013.03	2012.95
Education_Basic	0.06	0.04	0.00	0.00
Education_Graduation	0.51	0.48	0.52	0.53
Education_Master	0.15	0.19	0.16	0.14
Education_PhD	0.20	0.19	0.24	0.25
Marital_Status_Alone	0.01	0.00	0.00	0.00
Marital_Status_Divorced	0.00	0.13	0.08	0.12
Marital_Status_Married	0.00	0.51	0.42	0.37
Marital_Status_Single	0.98	0.00	0.20	0.19
Marital_Status_Together	0.00	0.33	0.25	0.26

<b>Marital_Status_Widow</b>	0.00	0.03	0.05	0.05
<b>Marital_Status_YOLO</b>	0.01	0.00	0.00	0.00

```
# Number of customers in each cluster
cluster_counts = data['Cluster'].value_counts().sort_index()
print("Number of customers per cluster:")
print(cluster_counts)
```

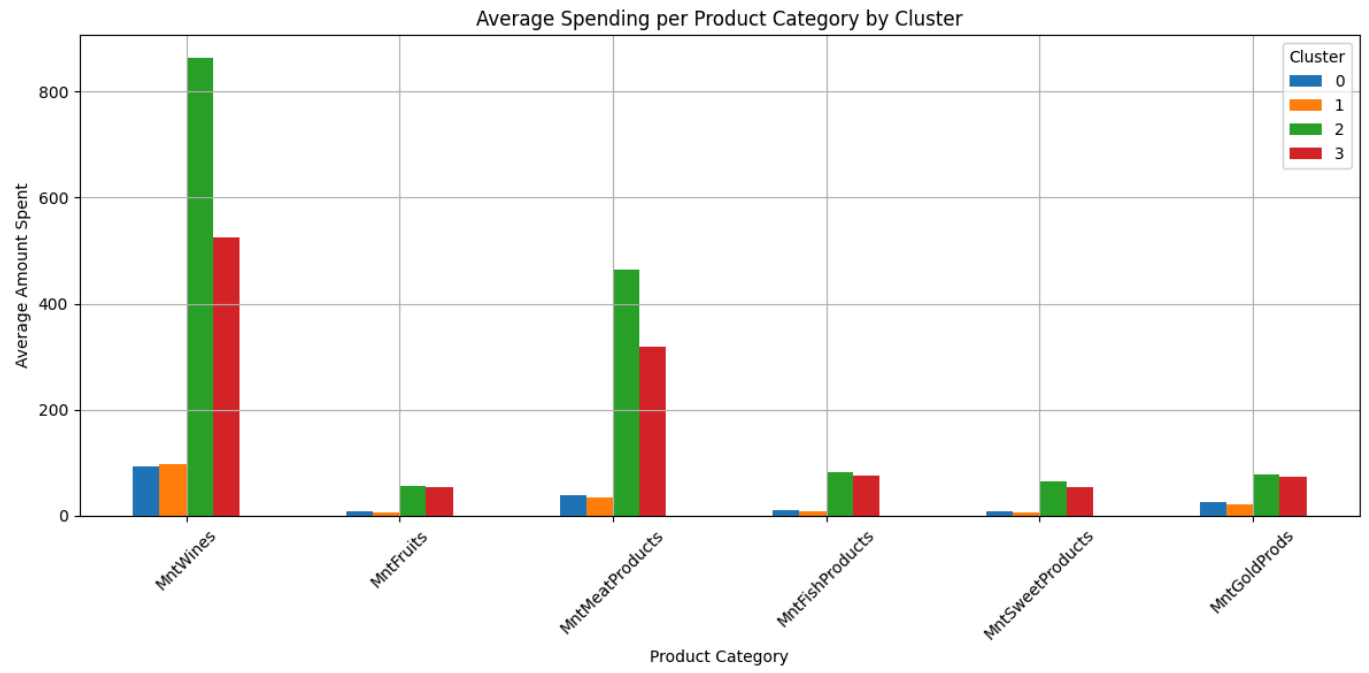
```
➦ Number of customers per cluster:
Cluster
0      294
1      993
2      195
3      734
Name: count, dtype: int64
```

```
# Visualize average spending per cluster
spending_features = [
    'MntWines', 'MntFruits', 'MntMeatProducts',
    'MntFishProducts', 'MntSweetProducts', 'MntGoldProds'
]

avg_spending = data.groupby('Cluster')[spending_features].mean()

plt.figure(figsize=(10, 6))
avg_spending.T.plot(kind='bar', figsize=(12, 6))
plt.title('Average Spending per Product Category by Cluster')
plt.xlabel('Product Category')
plt.ylabel('Average Amount Spent')
plt.xticks(rotation=45)
plt.legend(title='Cluster', loc='upper right')
plt.grid(True)
plt.tight_layout()
plt.show()
```

↔ <Figure size 1000x600 with 0 Axes>



```
# Comparing income distribution by cluster visual
plt.figure(figsize=(8, 5))
sns.boxplot(data=data, x='Cluster', y='Income', palette='Pastel1')
plt.title('Income Distribution by Cluster')
plt.ylim(0, 150000) # Adjustable the upper limit as needed
plt.grid(True)
plt.show()
```

↗ /tmp/ipython-input-149-3784089105.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign

```
sns.boxplot(data=data, x='Cluster', y='Income', palette='Pastel1')
```



#Cluster vs Web Activity

```
plt.figure(figsize=(8, 5))
```

```
sns.barplot(data=data, x='Cluster', y='NumWebPurchases', ci=None, palette='Set3')
```

```
plt.title('Average Web Purchases per Cluster')
```

```
plt.grid(True)
```

```
plt.show()
```

↗ /tmp/ipython-input-150-486649454.py:3: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=data, x='Cluster', y='NumWebPurchases', ci=None, palette='Set3')  
/tmp/ipython-input-150-486649454.py:3: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign

```
sns.barplot(data=data, x='Cluster', y='NumWebPurchases', ci=None, palette='Set3')
```

